

pdfmole

The **pdfmole** package provides a bundle of functions designed to assist in extracting tables from PDF-files. To read-in the data either

- **pdfminer**,
- **pdfboxr** or
- **tesseract** can be used.

In principle, any package which returns the data in a similar format could be used. The packages **pdfminer** and **pdfboxr** can be used if the PDF-file store already the text (in most cases) if the PDF contains only images of the tables **tesseract** can be used.

1. Installation

```
# install tesseract (optional)
install.packages("tesseract")

# install pdfminer (optional)
remotes::install_github("FlorianSchwendinger/pdfminer")

# install pdfboxr (optional)
remotes::install_gitlab("schwe/pdfboxr")

# install pdfmole
remotes::install_github("ben-schwen/pdfmole")
```

2. Workflow

```
options(width = 100)
library("pdfboxr")
library("pdfmole")
```

The basic work flow of **pdfmole** can be divided into

1. read data (`pdfboxr::read_chars()`)
2. align rows (`align_rows(x, method, ...)`)
3. aligns columns (`align_columns(x, method, ...)`)
4. group cells (`group_cells(x, collapse)`)
5. transform the data into a rectangular format (`mole(x, header, simplify, keep)`)

3. Example usage

3.1. Read data

Use **pdfboxr** (or **pdfminer**) to read the PDF-file into **R**,

```
pdf_file <- file.path(system.file("pdfs", package = "pdfmole"), "cars.pdf")
pdf <- pdfboxr::read_chars(pdf_file, pages = 1:2)
pdf
```

```
## A pdf document with 2 pages and
##   metainfo text fonts
## 1         2 313     3
## elements.
```

pdfboxr returns an object of class "pdf_document" which is a list containing several **data.frames**, the print method shows the number of rows of each **data.frame**. Each **data.frame** contains specific information about the PDF-file, for example the **data.frame** contains meta information about the PDF-file.

```
pdf$metainfo
```

```
##   rotation y0 x0  y1  x1 pid
## 1         0 0  0 842 595   1
## 2         0 0  0 842 595   2
```

The columns are defined as follows.

Name	Description
pid	page id
rotation	the rotation of the page
x0	the bounding box start coordinate on the x-axis
x1	the bounding box end coordinate on the x-axis
y0	the bounding box start coordinate on the y-axis
y1	the bounding box end coordinate on the y-axis

The **data.frame** named **text** contains the *page id*, *font name*, *font size*, *name of the color space*, *color* and *bounding box* for each character of the PDF-file.

```
d0 <- pdf$text
d0[c("rotation", "yscale", "xscale")] <- NULL # so the output fits to the column width
head(d0)
```

```
##   pid text                font size    x0    y0    x1    y1 width height
## 1   1   c BAAAAA+LiberationSans  10 288.20 769.12 293.20 774.7  5.00   5.58
## 2   1   a BAAAAA+LiberationSans  10 293.20 769.12 298.76 774.7  5.56   5.58
## 3   1   r BAAAAA+LiberationSans  10 298.79 769.12 302.12 774.7  3.33   5.58
## 4   1   s BAAAAA+LiberationSans  10 302.09 769.12 307.09 774.7  5.00   5.58
## 5   1   s          Courier-Bold  12  77.20 747.29  84.40 753.6  7.20   6.31
## 6   1   p          Courier-Bold  12  84.40 747.29  91.60 753.6  7.20   6.31
```

Based on this information we want to assign a row and column id to each row of the **data.frame**. This will allow us to transform character level **data.frame** into the typical format. Additionally we can use information like the location, font size or font type to filter out specific characters / blocks / rows.

3.2 Align blocks (optional)

The block information is assigned by **pdfminer** and signals which characters belong to the same text block. The function **group_blocks** can be used to paste together all the characters belonging to the same block. **pdfmoler** provides no block information therefore we skip this part.

```
d <- group_blocks(d0)
head(d, 3)
```

3.3 Align rows

The function `align_rows(x, method, ...)` currently implements three methods "exact_match", "hclust" and "fixed_width".

```
d <- align_rows(d0)
head(d, 30)
```

##	pid	text	font	size	x0	y0	x1	y1	width	height	row
## 1	1	c	BAAAAA+LiberationSans	10	288.20	769.12	293.20	774.7	5.00	5.58	1
## 2	1	a	BAAAAA+LiberationSans	10	293.20	769.12	298.76	774.7	5.56	5.58	1
## 3	1	r	BAAAAA+LiberationSans	10	298.79	769.12	302.12	774.7	3.33	5.58	1
## 4	1	s	BAAAAA+LiberationSans	10	302.09	769.12	307.09	774.7	5.00	5.58	1
## 5	1	s	Courier-Bold	12	77.20	747.29	84.40	753.6	7.20	6.31	2
## 6	1	p	Courier-Bold	12	84.40	747.29	91.60	753.6	7.20	6.31	2
## 7	1	e	Courier-Bold	12	91.60	747.29	98.80	753.6	7.20	6.31	2
## 8	1	e	Courier-Bold	12	98.80	747.29	106.00	753.6	7.20	6.31	2
## 9	1	d	Courier-Bold	12	106.00	747.29	113.20	753.6	7.20	6.31	2
## 10	1	d	Courier-Bold	12	124.40	747.29	131.60	753.6	7.20	6.31	2
## 11	1	i	Courier-Bold	12	131.60	747.29	138.80	753.6	7.20	6.31	2
## 12	1	s	Courier-Bold	12	138.80	747.29	146.00	753.6	7.20	6.31	2
## 13	1	t	Courier-Bold	12	146.00	747.29	153.20	753.6	7.20	6.31	2
## 14	1	1	Courier	12	67.90	732.27	75.10	738.6	7.20	6.33	4
## 15	1	4	Courier	12	115.20	732.27	122.40	738.6	7.20	6.33	4
## 16	1	2	Courier	12	154.80	732.27	162.00	738.6	7.20	6.33	4
## 17	1	2	Courier	12	67.90	717.27	75.10	723.6	7.20	6.33	5
## 18	1	4	Courier	12	115.20	717.27	122.40	723.6	7.20	6.33	5
## 19	1	1	Courier	12	147.60	717.27	154.80	723.6	7.20	6.33	5
## 20	1	0	Courier	12	154.80	717.27	162.00	723.6	7.20	6.33	5
## 21	1	3	Courier	12	67.90	702.27	75.10	708.6	7.20	6.33	6
## 22	1	7	Courier	12	115.20	702.27	122.40	708.6	7.20	6.33	6
## 23	1	4	Courier	12	154.80	702.27	162.00	708.6	7.20	6.33	6
## 24	1	4	Courier	12	67.90	687.27	75.10	693.6	7.20	6.33	7
## 25	1	7	Courier	12	115.20	687.27	122.40	693.6	7.20	6.33	7
## 26	1	2	Courier	12	147.60	687.27	154.80	693.6	7.20	6.33	7
## 27	1	2	Courier	12	154.80	687.27	162.00	693.6	7.20	6.33	7
## 28	1	5	Courier	12	67.90	672.27	75.10	678.6	7.20	6.33	8
## 29	1	8	Courier	12	115.20	672.27	122.40	678.6	7.20	6.33	8
## 30	1	1	Courier	12	147.60	672.27	154.80	678.6	7.20	6.33	8


After applying the function the column row is added to the `data.frame`.

3.3 Align columns

Visualization


Pixelplot

```
pixelplot(d, scale = 0.35, cex.axis = 0.35)
```



Bounding box plot

```
bboxplot(d, pid = 1L)
```



3.6 Group cells

```
d <- group_cells(d)
head(d, 3)
```

```
##   pid row col  text
## 1   1   2   2 speed
## 2   1   2   3  dist
## 3   1   4   1     1
```

3.7 Transform into a rectangular (dense) format

After the grouping the data is available similar to a simple triplet sparse matrix format (sometimes called coordinate format). To obtain the usual (dense) `data.frame` / `matrix` format the function `mole()` can be used.

```
x <- mole(d, header = TRUE, simplify = TRUE)
x
```

```
## X1    speed  dist
## 1     4      2
## 2     4     10
## 3     7      4
## 4     7     22
## 5     8     16
## 6     9     10
## ...   ...   ...
## 8     10     26
## 9     10     34
## 11    11     28
## 12    12     14
## 13    12     20
## 14    12     24
## 16    13     26
## 18    13     34
## 21    14     36
## 24    15     20
## 26    15     54
## 28    16     40
## 32    18     42
## 34    18     76
## 35    18     84
## 39    20     32
## ...   ...   ...
## 45    23     54
## 46    24     70
## 47    24     92
## 48    24     93
## 49    24    120
## 50    25     85
```

The function `mole()` returns an object of class `mole`,

```
str(x)
```

```
## List of 3
## $ X1    : int [1:50] 1 2 3 4 5 6 7 8 9 10 ...
```

```
## $ speed: int [1:50] 4 4 7 7 8 9 10 10 10 11 ...
## $ dist : int [1:50] 2 10 4 22 16 10 18 26 34 17 ...
## - attr(*, "class")= chr "mole"
```

which can be coerced to either `data.frame`

```
head(as.data.frame(x))
```

```
##   X1 speed dist
## 1  1     4    2
## 2  2     4   10
## 3  3     7    4
## 4  4     7   22
## 5  5     8   16
## 6  6     9   10
```

or `matrix`.

```
head(as.matrix(x))
```

```
##      X1 speed dist
## [1,]  1     4    2
## [2,]  2     4   10
## [3,]  3     7    4
## [4,]  4     7   22
## [5,]  5     8   16
## [6,]  6     9   10
```